PCA Analysis of RNA-seq Data Using DESeq2

Oğuzhan Işılay

2025-12-30

Introduction

RNA-seq experiments measure the expression levels of thousands of genes at the same time.
Because of this, the dataset becomes very large and complex. It is difficult to understand the relationships between samples by looking directly at the raw data.

One common way to explore this type of data is **Principal Component Analysis (PCA)**.

PCA is a **dimensionality reduction method**.
It simplifies complex datasets and helps us visualize patterns between samples.

In RNA-seq analysis, PCA is mainly used to:

- visualize similarity between samples
- detect unexpected samples (outliers)
- evaluate whether treatment and control groups separate

It is important to remember that **PCA is not a statistical test**.
It is mainly used for **exploratory analysis and visualization**.

When Can PCA Be Applied?

PCA cannot be applied to every dataset. Some conditions must be satisfied:

- The data must be **numerical**
- There must be **multiple samples**
- The data structure must be correct

In most biological datasets:

- **rows represent genes**
- **columns represent samples**

RNA-seq Data and PCA

Raw RNA-seq data consist of **read counts**.

However, count data have large differences in variance between genes. Some genes have very high counts while others have very low counts.

Because of this, raw counts are **not suitable for PCA directly**.

Before performing PCA, the data should be transformed using methods that stabilize variance.

One of the most common tools for this purpose is **DESeq2**.

What is DESeq2?

DESeq2 is a widely used **R/Bioconductor package** for RNA-seq analysis.

It is mainly used to identify **differentially expressed genes** between experimental conditions.

However, DESeq2 also provides useful transformations that make RNA-seq data suitable for visualization and exploratory analysis.

One of these transformations is called:

Variance Stabilizing Transformation (VST).

VST helps to:

- stabilize variance across genes
- reduce the effect of extremely high counts
- make samples more comparable

After applying VST, the data become more suitable for PCA.

Dataset Used in This Example

In this example we use the **airway dataset**.

This dataset is commonly used for teaching RNA-seq analysis because:

- it comes from a real RNA-seq experiment
- it is relatively small
- sample metadata is already included
- it contains clear treatment and control groups

This makes it ideal for learning the DESeq2 workflow.

Loading Required Packages

First we install and load the packages that we will use in this analysis.

```
install.packages("BiocManager") # Install Bioconductor manager (only needed once)
BiocManager::install(c("airway","DESeq2")) # Install the required bioinformatics packages
```

```
# Load the airway dataset package
library(airway)

# Load the DESeq2 package for RNA-seq analysis
library(DESeq2)
```

Loading the Dataset

The **airway dataset** contains both gene expression counts and sample metadata.

```
# Load the airway dataset
data(airway)

# Extract the count matrix
# Rows = genes
# Columns = samples
counts <- assay(airway)

# Extract metadata describing each sample
# For example: treatment condition, sample name, etc.
coldata <- colData(airway)
```

Creating a DESeq2 Dataset

To analyze RNA-seq data with **DESeq2**, we first create a **DESeqDataSet** object.

This object combines three important components:

- the **count matrix** (gene expression counts)
- the **sample metadata**
- the **experimental design**

```
dds <- DESeqDataSetFromMatrix(
  countData = counts,    # gene expression counts
  colData = coldata,     # sample information (metadata)
  design = ~ dex         # experimental condition
)
```

Here the variable dex represents the experimental condition (treatment vs control).

Running the DESeq2 Pipeline

Next, we run the **DESeq2 analysis pipeline**.

```
dds <- DESeq(dds) # Run the main DESeq2 analysis
```

This step performs several important operations:

- **normalization** – adjusts for differences in sequencing depth between samples
- **dispersion estimation** – estimates variability in gene expression
- **model fitting** – fits a statistical model for each gene

These steps prepare the data for downstream analyses such as PCA and differential expression analysis.

Variance Stabilizing Transformation

Before performing PCA, we transform the RNA-seq count data using **Variance Stabilizing Transformation (VST)**.

```
vsd <- vst(dds, blind = FALSE) # Apply variance stabilizing transformation
```

This transformation helps to:

- stabilize the variance across genes
- reduce the effect of extremely high counts
- make samples easier to compare

After this step, the data become more suitable for visualization methods such as **PCA**.

PCA Visualization

Now we can generate a **PCA plot**.

```
plotPCA(vsd, intgroup = "dex")
```

```
## using ntop=500 top features by variance
```

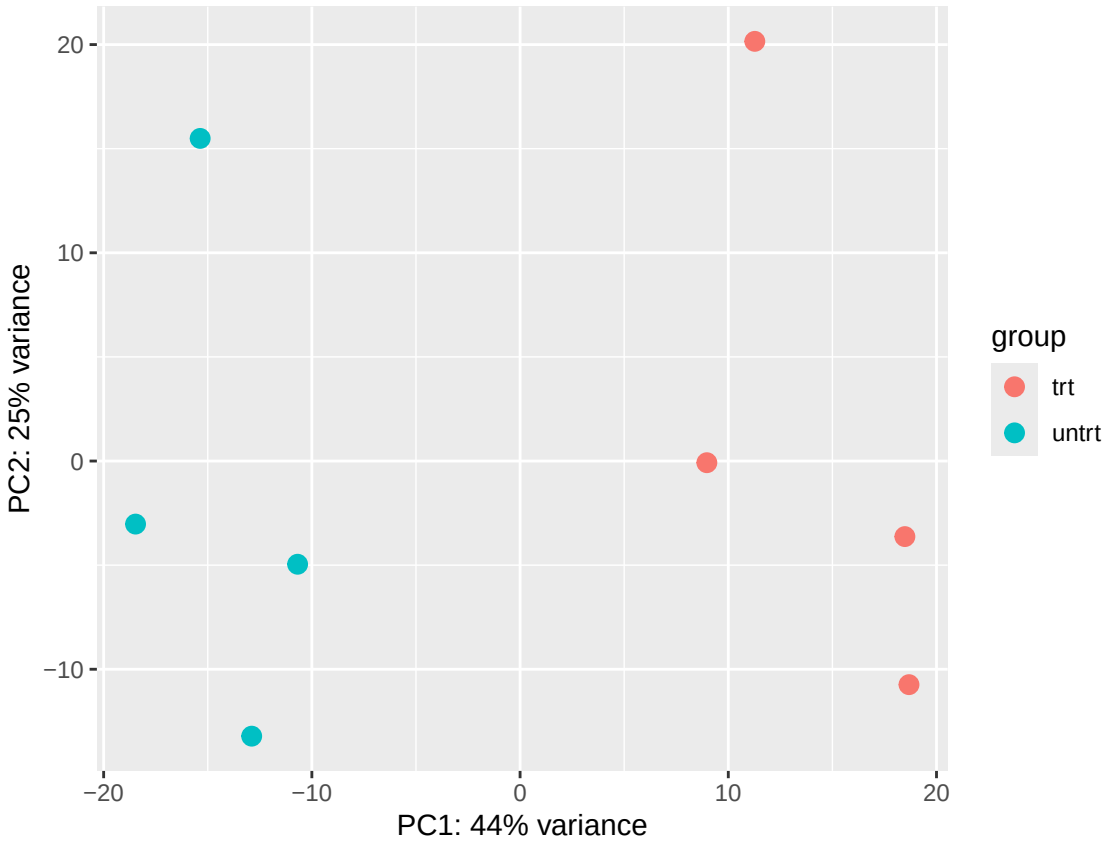

Each point in the plot represents **one sample**.

Samples that have similar gene expression patterns will appear close to each other in the plot.

How to Interpret the PCA Plot

The PCA plot contains two main axes:

PC1 (Principal Component 1)
Represents the largest source of variation in the dataset.

PC2 (Principal Component 2)
Represents the second largest source of variation.

Together, PC1 and PC2 summarize a large portion of the variability in the dataset.

In RNA-seq PCA plots we usually expect:

- samples from the same group to cluster together
- treatment and control samples to separate

If a sample appears far away from others, it may indicate a potential **outlier or technical issue**.

Summary

In this tutorial we learned:

- what PCA is
- why PCA is useful for RNA-seq data
- why raw count data must be transformed
- how DESeq2 and VST prepare the data for PCA
- how to visualize samples using PCA

PCA is an important **quality control step** in RNA-seq analysis and helps researchers understand the structure of their data before performing more advanced analyses.